

Zachary Cangemi

New York, NY | zcangemi@gmail.com
zacharycangemi.com | github.com/zacangemi | linkedin.com/in/zachary-cangemi-0a08b4b1

EXPERIENCE

Broadridge Financial Solutions, Inc.

Newark, NJ

Senior Data Scientist & AI Engineer

2024–Present

- Fine-tuned Meta Llama 3.1 8B across two RTX 3090 GPUs using Axolotl to create a specialized tax-domain model, preparing JSONL-formatted instruction data and testing model behavior locally.
- Work directly with two senior leaders in a small, autonomous AI function to evaluate open-weight models, build local inference workflows, and implement LLM solutions for financial-services use cases.
- Collaborate with Broadridge's Wealth AI group on enterprise AI strategy and the use of local LLMs for workloads containing personally identifiable information (PII) and other sensitive financial data.
- Built a high-volume, on-premises OCR pipeline using two local models to parse sensitive documents and forms, which became a client-serving product that kept data in-house.

Senior Data Scientist

2021–2024

- Owned product data workflows supporting tax reporting, including import mapping and backend data preparation, while partnering with engineering to define requirements and resolve complex data issues.
- Built Python and SQL tools for onboarding, reconciliation, data-quality analysis, dashboards, and product research across security-reference, transactional, and cost-basis datasets.

Associate Product Data Analyst

2019–2021

- Managed data onboarding for large institutional-bank clients collectively representing more than \$60 million in contract revenue, translating complex datasets into validated tax-reporting workflows.
- Built Python and SQL automation using pandas, psycopg2, and CSV-processing tools to streamline inbound and outbound file processing; also served as a primary contributor to a new platform for ingesting, enriching, transforming, and reporting brokerage data.

Data Analyst

2018–2019

- Built Python and SQL data-aggregation workflows to transform and onboard complex client datasets while diagnosing data and product issues directly with clients.
- Developed SQL and Python extraction and reconciliation scripts using multi-table joins, temporary tables, and validation checks to support accurate downstream processing and tax-reporting workflows.

EDUCATION

Colby College — Waterville, ME

2013–2017

Bachelor of Arts in Environmental Science, Concentration in Water Science

University of Otago — Dunedin, New Zealand

2016

SKILLS

Programming: Python, SQL

Inference Engineering: vLLM, llama.cpp, local model serving, hybrid CPU/GPU inference, tensor parallelism, quantized model deployment, speculative decoding, prefix caching, KV cache optimization, GPU performance profiling

Agent Systems: Agent workflows, tool calling, tool integration, task orchestration, scheduled automation, Hermes Agent, OpenCode

Model Development & Evaluation: Axolotl, LLM fine-tuning, instruction-data preparation, evaluation harnesses, model and agent capability testing, failure analysis, long-context evaluation, multimodal evaluation

Systems: Linux, Git, unit testing, multi-GPU systems

LANGUAGES

English (native), Portuguese (Brazilian; proficient), Spanish (proficient)